



**UNIVERSITY OF
WATERLOO**

**Midterm Examination
Fall 2017**

**Computer Science 343
Concurrent and Parallel Programming
Sections 001, 002**

**Duration of Exam: 1 hour 50 minutes
Number of Exam Pages (including cover sheet): 5
Total number of questions: 6
Total marks available: 102**

CLOSED BOOK, NO ADDITIONAL MATERIAL ALLOWED

**Instructor: Peter Buhr
November 1, 2017**

1. (a) i. **3 marks** Rewrite the following two **BAD** forms of loop exit into **GOOD** forms.

| BAD | BAD |
|--|--|
| <pre> for (;;) { S1 if (C1) { S2 } else { break; } S3 } </pre> | <pre> for (;;) { S1 if (C1) { break; } else { S2 } S3 } </pre> |

- ii. **1 mark** Explain the problem with the **BAD** form.
- (b) **2 marks** Why is a labelled **break** statement better than using a **goto** statement to perform the same operation?
- (c) **1 mark** What property makes a variable a *flag variable*?
- (d) i. **2 marks** State two problems with *return codes* to indicate multiple outcomes from a routine.
 ii. **1 mark** Are *return codes* faster or slower than *exceptions* for indicating multiple outcomes?
- (e) **2 marks** Explain how the **_Finally** clause works with respect to execution of a **try** statement.
- (f) **2 marks** Explain the terms *source* and *faulting* execution with respect to exception handling.
- (g) i. **1 mark** Where does control return at the end of a **catch** clause of a **try** statement?
 ii. **1 mark** Where does control return at the end of a **_CatchResume** clause of a **try** statement?
- (h) **1 mark** Why is **_Throw E() _At coroutine/task** unsupported in $\mu\text{C++}$?
- (i) **1 mark** State the primary criteria for using heap allocation.
- (j) **1 mark** Why does concurrent heap-allocation cause significant performance problems?
2. (a) **2 marks** Explain why coroutines are sequential versus concurrent.
- (b) **1 mark** What property of a $\mu\text{C++}$ coroutine allows modularization within the coroutine main?
- (c) **1 mark** When a coroutine's interface-member is called, on which stack does the call-frame go?
- (d) **2 marks** When a non-terminated coroutine is deallocated, what occurs to that coroutine's stack and why does it occur?
- (e) **1 mark** What does it mean to *linearize* (or *flatten*) a series of executable statements?
- (f) **1 mark** What is the purpose of the $\mu\text{C++}$ *verify* routine for a coroutine or task?
- (g) **1 mark** Given a full-coroutine cycle that is executing, what happens when suspend calls are performed?
- (h) **2 marks** Explain the difference between Python and $\mu\text{C++}$ coroutines. Give an example of what the difference does or does not allow.
3. (a) **1 mark** What is a concurrent *bottleneck*?
- (b) **1 mark** What is the most important take-away (lesson) from Amdahl's law?
- (c) **1 mark** What is the *critical path* in concurrent execution?
- (d) **2 marks** Is COBEGIN/COEND as expressive as START/WAIT? Explain.
- (e) **1 mark** Explain why **static** variables are dangerous in a **_Task** type.
- (f) **2 marks** Explain the difference between liveness (rule 4) and eventual progress (rule 5) in the mutual-exclusion game.

- (g) **2 marks** The following is the declare-intent solution for mutual exclusion:

```
me = WantIn;           // entry protocol
while ( you == WantIn ) {}
CriticalSection();     // critical section
me = DontWantIn;       // exit protocol
```

Explain what rule of the mutual-exclusion game is violated and how.

- (h) **2 marks** Given the Lock variable and an implementation of the test-and-set instruction:

```
int Lock = OPEN; // shared
```

```
int TestSet( int & b ) {
    // begin atomic
    int temp = b;
    b = CLOSED;
    // end atomic
    return temp;
}
```

use them to write the entry and exit protocols for mutual exclusion.

4. (a) **2 marks** What are *independent* and *dependent* critical sections?
 - (b) **1 mark** How many locks are needed for mutual exclusion?
 - (c) **1 mark** How many checks does a blocking lock make before blocking a task?
 - (d) **2 marks** Explain the difference between *avoidance* and *prevention*.
 - (e) **2 marks** Why does a synchronization lock's wait member take a mutex-lock parameter?
 - (f) **2 marks** Explain why a mutex lock cannot perform synchronization, and why a synchronization lock cannot perform mutual exclusion.
5. **19 marks** Write a *semi-coroutine* with the following public interface (you may only add a public destructor and private members):

```
_Coroutine Grammar {
    char ch;           // character passed by caller
    void main();       // YOU WRITE THIS ROUTINE
public:
    _Event Match {};   // characters form a valid string in the language
    _Event Error {};   // last character results in string not in the language
    void next( char c ) {
        ch = c;
        resume();
    }
};
```

which verifies a string of characters matches the language $(^+_n (XY)^+_n)^+_n$, where X cannot be $'($, $n \geq 1$, and the number of open/closing parentheses and *repeated* pairs of XY characters are equal; i.e., there are 1 or more opening parenthesis, followed by the same number of repeated pairs of any characters, followed by the same number of closing parenthesis, e.g.:

| valid strings | invalid strings |
|---------------|-----------------|
| (ab) | (ab |
| ((xyxy)) | ((xy)) |
| (((#@#@#@))) | ab) |
| ((www)) | (a) |
| ((())) | (((((|
| ((X)X))) | ((() |
| ()() | (PPP) |

After creation, the coroutine is resumed with a series of characters (one character at a time). The coroutine accepts characters until:

- the characters form a valid string in the language, and it then raises the exception `Grammar::Match` at the last resumer;
- the last character results in a string not in the language, it then raises the exception `Grammar::Error` at the last resumer.

After the coroutine raises a `Match` or `Error` exception, it must terminate; sending more characters to the coroutine after this point is undefined. (You may use multiple **return** statements in `Grammar::main`.)

Write **ONLY** `Grammar::main`, do **NOT** write a main program that uses it! **No documentation or error checking of any form is required.**

Note: Few marks will be given for a solution that does not take advantage of the capabilities of the coroutine, i.e., you must use the coroutine's ability to retain data and execution state.

6. **31 marks** Divide and conquer is a technique that can be applied to certain kinds of problems. These problems are characterized by the ability to subdivide the work across the data, such that the work can be performed independently on the data. In general, the work performed on each group of data is identical to the work that is performed on the data as a whole. What is important is that only termination synchronization is required to know the work is done; the partial results can then be processed further.

Write a **COMPLETE** $\mu\text{C++}$ program to *efficiently* check if any row of a matrix of size $N \times M$ contains at least 2 Schmilblicks. For example, in:

$$\begin{pmatrix} 1 & -1 & 3 & 4 & -1 \\ 2 & 1 & 4 & -1 & 6 \\ 3 & -1 & -1 & 6 & -1 \\ -1 & 6 & 7 & -1 & 1 \\ 4 & -1 & 6 & 1 & 8 \end{pmatrix}$$

the Schmilblick value is -1 , and rows 1, 3, 4 contain at least 2 Schmilblicks. The matrix is checked concurrently along its rows. Each checking task has the following interface (you may only add a public destructor and private members):

```

_Event Schmilblick {};           // concurrent exception
_Task Schmilblooks {             // check row of matrix
    ...                           // YOU ADD HERE
    void main();                 // YOU WRITE THIS ROUTINE
public:
    _Event Stop {};             // concurrent exception
    Schmilblooks(                // YOU WRITE THIS ROUTINE
        const int row[],        // one row of the matrix
        const int cols         // number of columns in row
        uBaseTask & pgmMain,    // contact when Schmilblooks found
        int schmilblick        // schmilblick value
    );
};

```

The program main reads from standard input the Schmilblick value and matrix dimensions ($N \times M$), declares any necessary matrix, arrays and variables, reads (from standard input) and prints (to standard output) the matrix, concurrently checks the matrix values in each row, and prints a message to standard output if two Schmilblooks are found in any matrix row. **No documentation or error checking of any form is required.**

As an optimization, each Schmilblooks task that finds a second Schmilblick raises the concurrent exception Schmilblick at the pgmMain and then returns, and when the program main receives this concurrent exception, it raises exception Schmilblooks::Stop at any non-deleted Schmilblooks tasks. When the concurrent Stop exception is propagated in a Schmilblooks task, it stops performing the Schmilblooks check and returns.

An example of the input for the program is:

```

-1 5 5          Schmilblick value and matrix dimensions

1 -1 3 4 -1     matrix values
2 1 4 -1 6
3 -1 1 6 -1
4 -1 6 1 8
-1 6 7 -1 1

```

(The phrases “Schmilblick value and matrix dimensions” and “matrix values” do not appear in the input.) In general, the input format is free form, meaning any amount of white space may separate the values.

Example outputs are:

| | |
|----------------------------------|--------------------------------|
| 1, -1, 3, 4, -1, original matrix | 1, 2, 3, 4, 5, original matrix |
| 2, 1, 4, -1, 6, | 2, 1, 4, 5, 6, |
| 3, -1, 1, 6, -1, | 3, 4, 1, 6, 7, |
| 4, -1, 6, 1, 8, | 4, 5, 6, 1, 8, |
| -1, 6, 7, -1, 1, | 5, 6, 7, 8, 1, |
| Schmilblooks found | Schmilblooks not found |

(The phrase “original matrix” does not appear in the output.) Note, the comma is a terminator not a separator.